

regression1

November 5, 2023

Write a program to perform regression tasks over given data using direct functions and evaluate its performance.

NAME ATHARVA BALPANDE ROLL NO 313002 PRN 22111143 Div:C1

Dataset url "https://www.kaggle.com/code/gantalaswetha/usa-housing-dataset-linear-regression/input?select=USA_Housing.csv"

```
[ ]: #importing necessary libraries for data manipulation and visualization
import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from tensorflow import keras
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.linear_model import LinearRegression
from math import sqrt
from sklearn.metrics import mean_squared_error, r2_score
```

```
[ ]: #importing the dataset
file = 'USA_Housing.csv'
data=pd.read_csv(file)
```

```
[ ]: data.head()
```

```
[ ]: 
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	

	Avg. Area Number of Bedrooms	Area Population	Price	\
0	4.09	23086.800503	1.059034e+06	
1	3.09	40173.072174	1.505891e+06	
2	5.13	36882.159400	1.058988e+06	
3	3.26	34310.242831	1.260617e+06	

```
4                                4.23      26354.109472  6.309435e+05
```

```
                                Address
0  208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1  188 Johnson Views Suite 079\nLake Kathleen, CA...
2  9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3                                USS Barnett\nFPO AP 44820
4                                USNS Raymond\nFPO AE 09386
```

```
[ ]: #checking errors/null values/wrong values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                              5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

```
[ ]: #summary of the data
data.describe()
```

```
[ ]:      Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  \
count      5000.000000      5000.000000      5000.000000
mean      68583.108984        5.977222        6.987792
std       10657.991214        0.991456        1.005833
min       17796.631190        2.644304        3.236194
25%       61480.562388        5.322283        6.299250
50%       68804.286404        5.970429        7.002902
75%       75783.338666        6.650808        7.665871
max       107701.748378        9.519088       10.759588
```

```
      Avg. Area Number of Bedrooms  Area Population      Price
count      5000.000000      5000.000000  5.000000e+03
mean           3.981330      36163.516039  1.232073e+06
std           1.234137       9925.650114  3.531176e+05
min           2.000000       172.610686  1.593866e+04
25%           3.140000      29403.928702  9.975771e+05
50%           4.050000      36199.406689  1.232669e+06
75%           4.490000      42861.290769  1.471210e+06
```

max 6.500000 69621.713378 2.469066e+06

```
[ ]: data.head()
```

```
[ ]: Avg. Area Income Avg. Area House Age Avg. Area Number of Rooms \
0 79545.458574 5.682861 7.009188
1 79248.642455 6.002900 6.730821
2 61287.067179 5.865890 8.512727
3 63345.240046 7.188236 5.586729
4 59982.197226 5.040555 7.839388
```

```
Avg. Area Number of Bedrooms Area Population Price \
0 4.09 23086.800503 1.059034e+06
1 3.09 40173.072174 1.505891e+06
2 5.13 36882.159400 1.058988e+06
3 3.26 34310.242831 1.260617e+06
4 4.23 26354.109472 6.309435e+05
```

```
Address
0 208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1 188 Johnson Views Suite 079\nLake Kathleen, CA...
2 9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3 USS Barnett\nFPO AP 44820
4 USNS Raymond\nFPO AE 09386
```

```
[ ]: # Putting feature variable to X
X = data[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of
Rooms', 'Avg. Area Number of Bedrooms', 'Area Population']]

# Putting response variable to y
y = data['Price']
```

```
[ ]: pre_process = preprocessing.StandardScaler()
X = pd.DataFrame(pre_process.fit_transform(X))
```

```
[ ]: X.head()
```

```
[ ]:
0 1.028660 -0.296927 0.021274 0.088062 -1.317599
1 1.000808 0.025902 -0.255506 -0.722301 0.403999
2 -0.684629 -0.112303 1.516243 0.930840 0.072410
3 -0.491499 1.221572 -1.393077 -0.584540 -0.186734
4 -0.807073 -0.944834 0.846742 0.201513 -0.988387
```

```
[ ]: y.head()
```

```
[ ]: 0    1.059034e+06
      1    1.505891e+06
      2    1.058988e+06
      3    1.260617e+06
      4    6.309435e+05
      Name: Price, dtype: float64
```

```
[ ]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7,
      ↪, test_size = 0.3, random_state=2)
```

```
[ ]: print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)
```

```
(3500, 5)
(1500, 5)
(3500,)
(1500,)
```

```
[ ]: lm=LinearRegression()
```

```
[ ]: # fit the model to the training data
      lm.fit(X_train, y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: # print the intercept
      print(lm.intercept_)
```

```
1231006.3190642651
```

```
[ ]: #Let's see the coefficient
      coeff_df = pd.DataFrame(lm.coef_, X_test.columns, columns=['Coefficient'])
      coeff_df
```

```
[ ]:      Coefficient
      0  229275.224032
      1  163592.487140
      2  120113.253351
      3   3015.847572
      4  150556.340296
```

```
[ ]: # Making predictions using the model
      y_pred = lm.predict(X_test)
```

```
[ ]: mse = mean_squared_error(y_test, y_pred)
r_squared = r2_score(y_test, y_pred)
print('Mean_Squared_Error :',mse)
print('r_square_value :',r_squared)
```

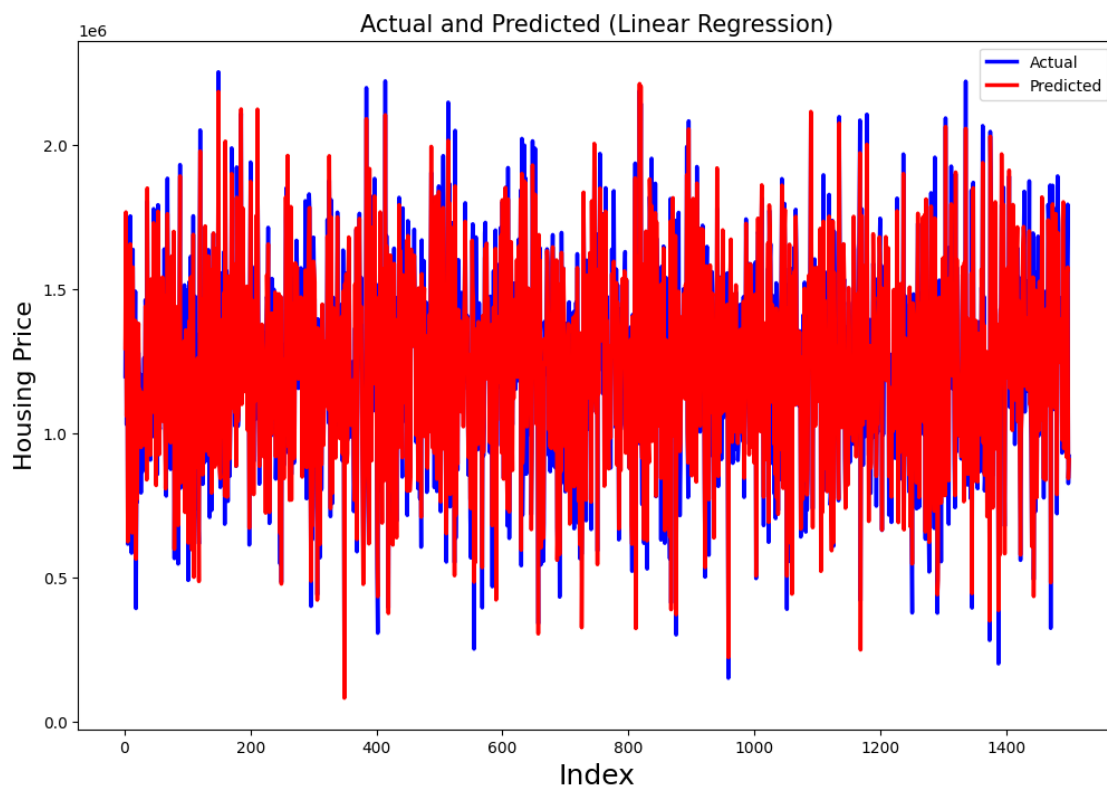
Mean_Squared_Error : 9831074697.740438
r_square_value : 0.9199287959786013

```
[ ]: c = [i for i in range(1, 1501, 1)] # Generating index

plt.figure(figsize=(12, 8))
plt.plot(c, y_test, color="blue", linewidth=2.5, linestyle="-", label="Actual")
    ↪ # Plotting Actual
plt.plot(c, y_pred, color="red", linewidth=2.5, linestyle="-",
    ↪ label="Predicted") # Plotting predicted

plt.title('Actual and Predicted (Linear Regression)', fontsize=15) # Plot
    ↪ heading
plt.xlabel('Index', fontsize=18) # X-label
plt.ylabel('Housing Price', fontsize=16)
plt.legend(loc="upper right") # Add a legend

plt.show()
```



[]: